

# Real Estate Price Prediction Model: Using Supervised Regression

## Load The Dataset

```
import pandas as pd

# Load the dataset
data = pd.read_csv('https://raw.githubusercontent.com/Jaden2802/Real_Estate_Prediction/main/House_Prices.csv')

# Display the first few rows
data.head()
```



	ID	Price	Bedrooms	Bathrooms	Sqft_living	Sqft_lot	Floors	Waterfront	View
0	1	280000.0	6	3.00	2400	9373	2.0	0	0
1	2	300000.0	6	3.00	2400	9373	2.0	0	0
2	3	647500.0	4	1.75	2060	26036	1.0	0	0
3	4	400000.0	3	1.00	1460	43000	1.0	0	0
4	5	235000.0	3	1.00	1430	7599	1.5	0	0

Next steps:

Generate code with data

 View recommended plots

## Data Exploration

```
# Get basic information about the dataset
data.info()

# Summary statistics
data.describe()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21613 entries, 0 to 21612
Data columns (total 20 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   ID                     21613 non-null  int64
1   Price                  21613 non-null  float64
2   Bedrooms               21613 non-null  int64
3   Bathrooms              21613 non-null  float64
4   Sqft_living            21613 non-null  int64
5   Sqft_lot               21613 non-null  int64
6   Floors                 21613 non-null  float64
7   Waterfront             21613 non-null  int64
8   View                   21613 non-null  int64
9   Condition              21613 non-null  int64
10  Grade                  21613 non-null  int64
11  Sqft_above             21613 non-null  int64
12  Sqft_basement          21613 non-null  int64
13  Yr_built               21613 non-null  int64
14  Yr_renovated           21613 non-null  int64
15  zipcode                21613 non-null  int64
16  Lat                    21613 non-null  float64
17  Long                   21613 non-null  float64
18  Sqft_living15          21613 non-null  int64
19  Sqft_lot15             21613 non-null  int64
dtypes: float64(5), int64(15)
memory usage: 3.3 MB
```

	ID	Price	Bedrooms	Bathrooms	Sqft_living	Sqft_lo
count	21613.00000	2.161300e+04	21613.000000	21613.000000	21613.000000	2.161300e+0
mean	10807.00000	5.401822e+05	3.370842	2.114757	2079.899736	1.510697e+0
std	6239.28002	3.673622e+05	0.930062	0.770163	918.440897	4.142051e+0
min	1.00000	7.500000e+04	0.000000	0.000000	290.000000	5.200000e+0
25%	5404.00000	3.219500e+05	3.000000	1.750000	1427.000000	5.040000e+0
50%	10807.00000	4.500000e+05	3.000000	2.250000	1910.000000	7.618000e+0
75%	16210.00000	6.450000e+05	4.000000	2.500000	2550.000000	1.068800e+0
max	21613.00000	7.700000e+06	33.000000	8.000000	13540.000000	1.651359e+0

▼ Data Cleaning

```
# Check for missing values
print(data.isnull().sum())

# Handle missing values (example: fill with median)
data = data.fillna(data.median())

# Remove duplicates if any
data = data.drop_duplicates()

#Remove unnecessary columns
#data = data.drop(columns=['ID', 'Sqft_living', 'Sqft_living15', 'Sqft_lot15', 'Grade', 'View', 'Sqft_lot', 'Sqft_above', 'Sqft_basement'])
```

```
ID      0
Price    0
Bedrooms 0
Bathrooms 0
Sqft_living 0
Sqft_lot  0
Floors    0
Waterfront 0
View      0
Condition 0
Grade     0
Sqft_above 0
Sqft_basement 0
Yr_built  0
Yr_renovated 0
zipcode   0
Lat       0
Long      0
Sqft_living15 0
Sqft_lot15  0
dtype: int64
```

## ✓ Exploratory Data Analysis (EDA)

```
import seaborn as sns
import matplotlib.pyplot as plt

#Distribution of target variable (price)
sns.histplot(data['Price'])
plt.show()

#Boxplot for each feature
plt.figure(figsize=(15,8))
sns.boxplot(x="Bathrooms", y="Price", data=data)

plt.figure(figsize=(15,8))
sns.boxplot(x="Bedrooms", y="Price", data=data)

plt.figure(figsize=(15,8))
sns.boxplot(x="Floors", y="Price", data=data)

#ScatterPlot
plt.figure(figsize=(15,8))
sns.scatterplot(x="Floors", y="Price", data=data)
#CountPlot
plt.figure(figsize=(15,8))
sns.countplot(x="Bathrooms", data=data)




plt.figure(figsize=(15,8))
sns.countplot(x="Bedrooms", data=data)




plt.figure(figsize=(15,8))
sns.countplot(x="Floors", data=data)


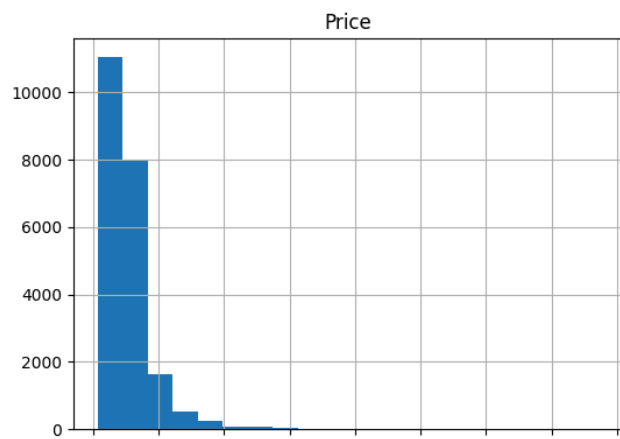



features = ['Price', 'Bedrooms', 'Bathrooms', 'Floors']
data[features].hist(bins=20, figsize=(14, 10))
plt.show()



#Correlation heatmap
corr_matrix = data.corr()
plt.figure(figsize=(15, 10))

sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', fmt='.2f')
plt.show()
```









ID	1.00	-0.02	0.00	0.00	-0.01	-0.14	0.02	-0.00	0.01	-0.02	0.01	-0.01	-0.00	0.02	-0.02	-0.01	-0.00	0.02	-0.00	-0.14
Price	-0.02	1.00	0.31	0.53	0.70	0.09	0.26	0.27	0.40	0.04	0.67	0.61	0.32	0.05	0.13	-0.05	0.31	0.02	0.59	0.08
Bedrooms	0.00	0.31	1.00	0.52	0.58	0.03	0.18	-0.01	0.08	0.03	0.36	0.48	0.30	0.15	0.02	-0.15	-0.01	0.13	0.39	0.03
Bathrooms	0.00	0.53	0.52	1.00	0.75	0.09	0.50	0.06	0.19	-0.12	0.66	0.69	0.28	0.51	0.05	-0.20	0.02	0.22	0.57	0.09
Sqft_living	-0.01	0.70	0.58	0.75	1.00	0.17	0.35	0.10	0.28	-0.06	0.76	0.88	0.44	0.32	0.06	-0.20	0.05	0.24	0.76	0.18
Sqft_lot	-0.14	0.09	0.03	0.09	0.17	1.00	-0.01	0.02	0.07	-0.01	0.11	0.18	0.02	0.05	0.01	-0.13	-0.09	0.23	0.14	0.72
Floors	0.02	0.26	0.18	0.50	0.35	-0.01	1.00	0.02	0.03	-0.26	0.46	0.52	-0.25	0.49	0.01	-0.06	0.05	0.13	0.28	-0.01
Waterfront	-0.00	0.27	-0.01	0.06	0.10	0.02	0.02	1.00	0.40	0.02	0.08	0.07	0.08	-0.03	0.09	0.03	-0.01	-0.04	0.09	0.03
View	0.01	0.40	0.08	0.19	0.28	0.07	0.03	0.40	1.00	0.05	0.25	0.17	0.28	-0.05	0.10	0.08	0.01	-0.08	0.28	0.07
Condition	-0.02	0.04	0.03	-0.12	-0.06	-0.01	-0.26	0.02	0.05	1.00	-0.14	-0.16	0.17	-0.36	-0.06	0.00	-0.01	-0.11	-0.09	-0.00
Grade	0.01	0.67	0.36	0.66	0.76	0.11	0.46	0.08	0.25	-0.14	1.00	0.76	0.17	0.45	0.01	-0.18	0.11	0.20	0.71	0.12
Sqft_above	-0.01	0.61	0.48	0.69	0.88	0.18	0.52	0.07	0.17	-0.16	0.76	1.00	-0.05	0.42	0.02	-0.26	-0.00	0.34	0.73	0.19
Sqft_basement	-0.00	0.32	0.30	0.28	0.44	0.02	-0.25	0.08	0.28	0.17	0.17	-0.05	1.00	-0.13	0.07	0.07	0.11	-0.14	0.20	0.02
Yr_built	0.02	0.05	0.15	0.51	0.32	0.05	0.49	-0.03	-0.05	-0.36	0.45	0.42	-0.13	1.00	-0.22	-0.35	-0.15	0.41	0.33	0.07
Yr_renovated	-0.02	0.13	0.02	0.05	0.06	0.01	0.01	0.09	0.10	-0.06	0.01	0.02	0.07	-0.22	1.00	0.06	0.03	-0.07	-0.00	0.01
zipcode	-0.01	-0.05	-0.15	-0.20	-0.20	-0.13	-0.06	0.03	0.08	0.00	-0.18	-0.26	0.07	-0.35	0.06	1.00	0.27	-0.56	-0.28	-0.15
Lat	-0.00	0.31	-0.01	0.02	0.05	-0.09	0.05	-0.01	0.01	-0.01	0.11	-0.00	0.11	-0.15	0.03	0.27	1.00	-0.14	0.05	-0.09
Long	0.02	0.02	0.13	0.22	0.24	0.23	0.13	-0.04	-0.08	-0.11	0.20	0.34	-0.14	0.41	-0.07	-0.56	-0.14	1.00	0.33	0.25
Sqft_living15	-0.00	0.59	0.39	0.57	0.76	0.14	0.28	0.09	0.28	-0.09	0.71	0.73	0.20	0.33	-0.00	-0.28	0.05	0.33	1.00	0.18
Sqft_lot15	-0.14	0.08	0.03	0.09	0.18	0.72	-0.01	0.03	0.07	-0.00	0.12	0.19	0.02	0.07	0.01	-0.15	-0.09	0.25	0.18	1.00

Double-click (or enter) to edit

## ✓ Split the Data

```
from sklearn.model_selection import train_test_split

# Define features and target variable
X = data.drop('Price', axis=1)
y = data['Price']

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=42)
```

## ✓ Model Training

```
import numpy as np

from sklearn.model_selection import cross_val_score
def classify(model, x, y):
    x_train, x_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=42)
    model.fit(x_train, y_train)
    print("Accuracy is", model.score(x_test, y_test)*100)
    # cross validation - it is used for better validation of model
    # eg: cv=5, train=4, test=1
    score = cross_val_score(model, x, y, cv=5)
    print("Cross validation is", np.mean(score)*100)
```

## Model Evaluation

## ✓ Linear Regression



```
from sklearn.linear_model import LinearRegression
model = LinearRegression()
classify(model, X, y)
```

↗ Accuracy is 70.3725085982858  
Cross validation is 69.57368511239943

## ▼ Decision Tree

```
from sklearn.tree import DecisionTreeRegressor
model = DecisionTreeRegressor()
classify(model, X, y)
```

↗ Accuracy is 73.25108202260915  
Cross validation is 71.17465615604863

## ▼ Random Forest

```
from sklearn.ensemble import RandomForestRegressor,ExtraTreesRegressor
model = RandomForestRegressor()
classify(model, X, y)
```

↗ Accuracy is 87.91164749695032  
Cross validation is 86.45652778192785

## ▼ Gradient Boosting

```
from sklearn.ensemble import GradientBoostingRegressor
model = GradientBoostingRegressor()
classify(model, X, y)
```

↗ Accuracy is 87.03179210149418  
Cross validation is 86.14383759131867